

Lab exam 04

SET – 9

Q1. A social media app needs to compute trending topics from millions of posts.

- a) Write AI-assisted optimized SQL query.
- b) Explain indexing and partitioning strategy.

Prompt :

Prompt for Q1(a) – AI-Assisted Optimized SQL Query

Prompt:

“Generate an AI-optimized SQL query to find trending topics from millions of social media posts. The query should efficiently count hashtags, calculate engagement, and return the most trending topics using best-practice SQL optimizations.”

Prompt for Q1(b) – Indexing & Partitioning Strategy

Prompt:

“Explain a simple and effective indexing and partitioning strategy that improves performance when analyzing trending topics from a very large social media posts table.

Code:

```
-- Schema for trending topics demo (Postgres)

-- posts table (partitioning example shown separately below)
CREATE TABLE IF NOT EXISTS posts (
    id BIGSERIAL PRIMARY KEY,
    user_id BIGINT,
    content TEXT,
    created_at TIMESTAMPTZ NOT NULL DEFAULT now(),
    likes INT DEFAULT 0,
    shares INT DEFAULT 0
);

-- topics table
CREATE TABLE IF NOT EXISTS topics (
    id BIGSERIAL PRIMARY KEY,
    name TEXT NOT NULL UNIQUE
);

-- mapping table linking posts to topics (hashtags or normalized topics)
CREATE TABLE IF NOT EXISTS post_topics (
    post_id BIGINT NOT NULL REFERENCES posts(id) ON DELETE CASCADE,
    topic_id BIGINT NOT NULL REFERENCES topics(id) ON DELETE CASCADE,
    PRIMARY KEY (post_id, topic_id)
);

-- Materialized view for scores (optional)
CREATE MATERIALIZED VIEW IF NOT EXISTS mv_topic_scores AS
WITH recent_posts AS (
    SELECT id, created_at, COALESCE(likes,0) AS likes, COALESCE(shares,0) AS shares
    FROM posts
    WHERE created_at >= now() - INTERVAL '24 hours'
)
SELECT pt.topic_id, t.name AS topic,
    SUM(
        EXP( - (EXTRACT(EPOCH FROM now() - recent_posts.created_at) / 3600) / 6 )
        * (1.0 + LN(1.0 + recent_posts.likes) + 2.0 * LN(1.0 + recent_posts.shares))
    ) AS score,
```

```
 README.md > abc # Trending Topics Demo
1  # Trending Topics Demo
5  - A Python script `src/trending.py` that provides:
6    - `etl` command to extract hashtags from `posts.content` into `topics` and `post_topics`.
7    - `compute` command to run an optimized trending-topics query with recency decay and
     engagement weighting.
8
9  Prerequisites
10 - PostgreSQL accessible via environment variables: `PGHOST`, `PGPORT`, `PGDATABASE`,
   `PGUSER`, `PGPASSWORD`.
11 - Python 3.8+.
12
13  Quick start
14  1. Install dependencies (PowerShell):
15
16  < ``powershell
17  python -m pip install -r requirements.txt
18  >
19
20  2. Apply schema to your database (using `psql`):
21
22  < ``powershell
23  psql "host=$env:PGHOST port=$env:PGPORT dbname=$env:PGDATABASE user=$env:PGUSER
   password=$env:PGPASSWORD" -f "sql/schema.sql"
24  >
25
26  3. (Optional) Seed `posts` with sample content, or insert real data.
27
28  4. Run the ETL to extract hashtags into `post_topics`:
29
30  < ``powershell
31  python src/trending.py etl --batch 1000
32  >
33
34  5. Compute top topics:
35
36  < ``powershell
37  python src/trending.py compute --window "24 hours" --half-life 6 --limit 50
```



```
51
1 #!/usr/bin/env python3
2 """Minimal demo: ETL hashtags -> post_topics and compute trending topics.
3
4 Usage:
5     pip install -r requirements.txt
6     set env vars: PGHOST, PGPORT, PGDATABASE, PGUSER, PGPASSWORD
7     python src/trending.py compute --window "24 hours" --half-life 6 --limit 50
8     python src/trending.py etl --batch 1000
9
10 This is a simple, synchronous demo. For production, use streaming ingestion and incremental ag-
11 """
12 import argparse
13 import os
14 import re
15 import sys
16 from typing import List, Tuple
17
18 import psycopg2
19 from psycopg2.extras import execute_values
20
21 HASHTAG_RE = re.compile(r"#([A-Za-z0-9_]+)")
22
23 def get_conn():
24     params = {
25         'host': os.getenv('PGHOST', 'localhost'),
26         'port': int(os.getenv('PGPORT', 5432)),
27         'dbname': os.getenv('PGDATABASE', 'postgres'),
28         'user': os.getenv('PGUSER', 'postgres'),
29         'password': os.getenv('PGPASSWORD', ''),
30     }
31     return psycopg2.connect(**params)
32
33 def etl_extract_hashtags(batch: int = 1000):
34     """Scan posts in batches, extract hashtags and populate topics + post_topics.
35     This naive ETL is fine for a one-time migration/demo. For production, perform
36     extraction at write-time or via a streaming pipeline.
37     """
38
```

```
33     def etl_extract_hashtags(batch: int = 1000):
34         conn = get_conn()
35         try:
36             with conn:
37                 with conn.cursor() as cur:
38                     last_id = 0
39                     while True:
40                         cur.execute(
41                             "SELECT id, content FROM posts WHERE id > %s ORDER BY id ASC LIMIT %s",
42                             (last_id, batch),
43                         )
44                         rows = cur.fetchall()
45                         if not rows:
46                             print('ETL complete.')
47                             break
48                         topic_inserts = [] # (name,)
49                         post_topic_pairs = [] # (post_id, topic_name)
50                         for post_id, content in rows:
51                             last_id = post_id
52                             if not content:
53                                 continue
54                             tags = set(m.group(1).lower() for m in HASHTAG_RE.finditer(content))
55                             for t in tags:
56                                 topic_inserts.append((t,))
57                                 post_topic_pairs.append((post_id, t))
58                         if not topic_inserts:
59                             continue
60                         # Upsert topics and collect ids
61                         # Use temporary table to bulk upsert names and then join back
62                         cur.execute("CREATE TEMP TABLE tmp_topics(name text) ON COMMIT DROP;")
63                         execute_values(cur, "INSERT INTO tmp_topics(name) VALUES %s", topic_inserts)
64                         cur.execute(
65                             "INSERT INTO topics (name) SELECT name FROM tmp_topics ON CONFLICT (name) DO UPDATE SET name = EXCLUDED.name"
66                         )
67                         # Get mapping name -> id
68                         cur.execute(
69                             "SELECT id, name FROM topics WHERE name IN (SELECT name FROM tmp_topics)"
70                         )
71                         topic_id_map = {name: id for id, name in cur.fetchall()}
72                         for post_id, topic in post_topic_pairs:
73                             cur.execute("UPDATE posts SET topic_id = %s WHERE id = %s", (topic_id_map[topic], post_id))
```

```
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
```

```
        )
    mapping = {name: tid for tid, name in cur.fetchall()}
    # Prepare post_topics bulk insert (use ids)
    post_topic_values: List[Tuple[int,int]] = []
    for post_id, name in post_topic_pairs:
        tid = mapping.get(name)
        if tid:
            post_topic_values.append((post_id, tid))
    if post_topic_values:
        # INSERT ... ON CONFLICT DO NOTHING to avoid duplicates
        execute_values(
            cur,
            "INSERT INTO post_topics(post_id, topic_id) VALUES %s ON CONFLICT
post_topic_values,
        )
    print(f"Processed up to post id {last_id}; inserted {len(post_topic_value
finally:
    conn.close()

def compute_trending(window: str = '24 hours', half_life_hours: float = 6.0, limit: int = 50):
    conn = get_conn()
    try:
        with conn:
            with conn.cursor() as cur:
                sql = """
WITH recent_posts AS (
    SELECT id, created_at, COALESCE(likes,0) AS likes, COALESCE(shares,0) AS shares
    FROM posts
    WHERE created_at >= now() - %s::interval
)
SELECT t.id AS topic_id, t.name AS topic, COUNT(*) AS mentions,
       SUM(
           EXP( - (EXTRACT(EPOCH FROM now() - rp.created_at) / 3600) / %s )
           * (1.0 + LN(1.0 + rp.likes) + 2.0 * LN(1.0 + rp.shares))
       ) AS score
    FROM post_topics pt
```

```

110     JOIN recent_posts rp ON rp.id = pt.post_id
111     JOIN topics t ON t.id = pt.topic_id
112     GROUP BY t.id, t.name
113     ORDER BY score DESC
114     LIMIT %s;
115     """
116     cur.execute(sql, (window, float(half_life_hours), int(limit)))
117     rows = cur.fetchall()
118     print(f"Top {limit} topics (window={window}, half_life={half_life_hours}h):")
119     print("rank\ttopic_id\tmentions\tscore\tname")
120     for i, (tid, name, mentions, score) in enumerate(rows, start=1):
121         print(f"{i}\t{tid}\t{mentions}\t{score:.4f}\t{name}")
122     return rows
123 finally:
124     conn.close()
125
126 def main(argv=None):
127     parser = argparse.ArgumentParser()
128     sub = parser.add_subparsers(dest='cmd')
129     p_etl = sub.add_parser('etl', help='Extract hashtags from posts into topics/post_topics')
130     p_etl.add_argument('--batch', type=int, default=1000)
131     p_compute = sub.add_parser('compute', help='Compute trending topics')
132     p_compute.add_argument('--window', default='24 hours')
133     p_compute.add_argument('--half-life', type=float, default=6.0)
134     p_compute.add_argument('--limit', type=int, default=50)
135     args = parser.parse_args(argv)
136     if args.cmd == 'etl':
137         etl_extract_hashtags(batch=args.batch)
138     elif args.cmd == 'compute':
139         compute_trending(window=args.window, half_life_hours=args.half_life, limit=args.limit)
140     else:
141         parser.print_help()
142
143 if __name__ == '__main__':
144     main()

```

Output:

Rank	Topic	Mentions	Score
1	ai	4	18.5162
2	ml	2	9.1416
3	breaking	1	8.9498
4	fun	2	5.1896
5	tutorial	1	4.4055
6	music	1	4.2776
7	future	1	4.0411
8	lol	1	0.9121
9	systems	1	0.7891

Prompt for Q2(a) – Sample Stored Procedure

Prompt:

“Write a simple stored procedure that calculates trending topics from a posts table and stores the results in a separate trending table.”

Prompt for Q2(b) – Performance Benefits**Prompt:**

“Explain the performance benefits of using stored procedures instead of running raw SQL queries repeatedly.”

Observation:

The trending analysis shows that ‘ai’ is the most dominant topic, achieving the highest number of mentions and a significantly higher trend score compared to all other topics, indicating strong user interest and engagement. The topics ‘ml’ and ‘breaking’ follow, showing moderate popularity driven by fewer mentions but relatively high scores, suggesting timely relevance. Mid-range topics such as ‘fun’, ‘tutorial’, ‘music’, and ‘future’ display decent interaction but do not reach the engagement levels seen in the top categories. Meanwhile, lower-scoring topics like ‘lol’ and ‘systems’ reveal minimal traction, indicating less user attention within the dataset. Overall, the distribution suggests a clear concentration of interest around technology-related topics, especially AI and machine learning, while other categories remain comparatively less discussed.

Q2. AI suggests converting SQL into stored procedures.

- a) Develop sample procedure.**
- b) State performance benefits**

Prompt:

“Write a simple stored procedure that calculates trending topics from a posts table and stores the results in a separate trending table.”

Prompt:

“Explain the performance benefits of using stored procedures instead of running raw SQL queries repeatedly.”

Observation:

1. The stored procedure makes the trending topic calculation more organized and automated.
2. It reduces the need to repeatedly run long SQL queries manually.
3. Processing happens inside the database, which improves speed and efficiency.
4. It ensures consistent results every time the procedure is executed.
5. Overall, it enhances performance and simplifies the workflow for large datasets.

Code:

```
"""
Demo: call stored procedures `usp_GetCustomerOrders` and `usp_AddOrder` using pyodbc.

Before running:
- Install requirements: `pip install -r requirements.txt`
- Update `conn_str` to point to your SQL Server and database.

This script demonstrates pagination and reading an OUTPUT parameter.
"""

import pyodbc
from datetime import datetime
import os

# Build a connection string automatically if the user hasn't configured one.
# Edit the `SERVER` and `DATABASE` values below for your environment if needed.
SERVER = os.environ.get('SQL_SERVER', '(localdb)\\MSSQLLocalDB')
DATABASE = os.environ.get('SQL_DATABASE', 'master')

# Prefer newer Microsoft drivers if present, fall back to older 'SQL Server'
available_drivers = [d for d in pyodbc.drivers()]
preferred = ['ODBC Driver 18 for SQL Server', 'ODBC Driver 17 for SQL Server', 'SQL Server']
driver = next((d for d in preferred if d in available_drivers), None)
if not driver:
    raise SystemExit('No suitable ODBC driver found. Install Microsoft ODBC Driver for SQL Server')

conn_str = f"DRIVER={{{driver}}};SERVER={SERVER};DATABASE={DATABASE};Trusted_Connection=yes"

print('Using connection string:', conn_str)

def get_customer_orders(customer_id, page=1, page_size=5, start_date=None, end_date=None):
    with pyodbc.connect(conn_str) as conn:
        cur = conn.cursor()
        # Call the stored procedure; it returns rows that include TotalCount via COUNT() OVER()
        cur.execute("EXEC dbo.usp_GetCustomerOrders ?, ?, ?, ?, ?", customer_id, start_date, end_date, None, None)
        rows = cur.fetchall()
        if not rows:
```

```
    return [], 0
total = rows[0].TotalCount if hasattr(rows[0], 'TotalCount') else rows[0][5]
# map rows to dict for readability
orders = [
{
    'OrderId': r.OrderId,
    'CustomerId': r.CustomerId,
    'OrderDate': r.OrderDate,
    'Amount': float(r.Amount),
    'Status': r.Status,
}
for r in rows
]
return orders, total

def add_order(customer_id, amount, status='Processing', order_date=None):
    order_date = order_date or datetime.now()
    with pyodbc.connect(conn_str) as conn:
        cur = conn.cursor()
        # Use a T-SQL block to capture OUTPUT parameter
        cur.execute(
            "DECLARE @NewOrderId INT; EXEC dbo.usp_AddOrder ?, ?, ?, ?, @NewOrderId OUTPUT; SE
             customer_id, order_date, amount, status
        )
        new_id_row = cur.fetchone()
        new_id = new_id_row[0] if new_id_row else None
    return new_id

def main():
    customer_id = 1
    print('Adding a new order for customer', customer_id)
    new_id = add_order(customer_id, 199.99, status='Processing')
    print('New OrderId:', new_id)

    print('\nFetching page 1 of orders for customer', customer_id)
```

```
print('\nFetching page 1 of orders for customer', customer_id)
orders, total = get_customer_orders(customer_id, page=1, page_size=5)
print('Total orders for customer:', total)
for o in orders:
    print(o)

def mock_demo():
    # Mocked data path when a DB is not available – prints the same shape as the real demo
    print('Running mock demo (no DB connection)')
    customer_id = 1
    # Simulate adding an order
    new_id = 101
    print('Adding a new order for customer', customer_id)
    print('New OrderId:', new_id)

    # Simulate fetching paged results
    sample_orders = [
        {'OrderId': 101, 'CustomerId': 1, 'OrderDate': '2025-10-20', 'Amount': 199.99, 'Status': 'P'},
        {'OrderId': 3, 'CustomerId': 1, 'OrderDate': '2025-10-01', 'Amount': 250.00, 'Status': 'P'},
        {'OrderId': 2, 'CustomerId': 1, 'OrderDate': '2025-09-25', 'Amount': 45.00, 'Status': 'P'}
    ]
    total = 3
    print('\nFetching page 1 of orders for customer', customer_id)
    print('Total orders for customer:', total)
    for o in sample_orders:
        print(o)

if __name__ == '__main__':
    try:
        main()
    except Exception as exc:
        # If it's a DB connectivity error, fall back to mock demo so the script still demonstrates
        import pyodbc
        if isinstance(exc, pyodbc.Error) or 'SQL' in str(exc) or 'ODBC' in str(exc):
            print('Database connection failed, falling back to mock demo')
```

```

✓ if __name__ == '__main__':
✓   try:
✓     | main()
✓   except Exception as exc:
✓     | # If it's a DB connectivity error, fall back to mock demo so the script still demonstr
✓     | import pyodbc
✓     | if isinstance(exc, pyodbc.Error) or 'SQL' in str(exc) or 'ODBC' in str(exc):
✓     |   print('\nDatabase error detected:', exc)
✓     |   print('Falling back to mock demo output so you can see expected behavior.')
✓     |   mock_demo()
✓   else:
✓     | raise
| >  create_schema_and_data.sql
1  -- Creates schema and inserts sample data. Run in your SQL Server (adjust DB if need
2  IF OBJECT_ID('dbo.Orders', 'U') IS NOT NULL DROP TABLE dbo.Orders;
3  IF OBJECT_ID('dbo.Customers', 'U') IS NOT NULL DROP TABLE dbo.Customers;
4
5  CREATE TABLE dbo.Customers (
6    CustomerId INT IDENTITY(1,1) PRIMARY KEY,
7    Name NVARCHAR(100) NOT NULL,
8    Email NVARCHAR(100) NULL
9  );
10
11 CREATE TABLE dbo.Orders (
12   OrderId INT IDENTITY(1,1) PRIMARY KEY,
13   CustomerId INT NOT NULL FOREIGN KEY REFERENCES dbo.Customers(CustomerId),
14   OrderDate DATETIME NOT NULL DEFAULT GETDATE(),
15   Amount DECIMAL(10,2) NOT NULL,
16   Status NVARCHAR(20) NOT NULL
17 );
18
19 -- Insert sample customers
20 INSERT INTO dbo.Customers (Name, Email) VALUES
21 ('Alice Johnson', 'alice@example.com'),
22 ('Bob Smith', 'bob@example.com'),
23 ('Carol Lee', 'carol@example.com');
24
25 -- Insert sample orders
26 INSERT INTO dbo.Orders (CustomerId, OrderDate, Amount, Status) VALUES
27 (1, DATEADD(day,-30,GETDATE()), 120.50, 'Completed'),
28 (1, DATEADD(day,-25,GETDATE()), 45.00, 'Shipped'),
29 (1, DATEADD(day, -5, GETDATE()), 250.00, 'Processing'),
30 (2, DATEADD(day,-40,GETDATE()), 9.99, 'Completed'),
31 (2, DATEADD(day, -2, GETDATE()), 15.00, 'Processing'),
32 (3, DATEADD(day, -10, GETDATE()), 75.00, 'Completed');
33
34 PRINT 'Schema and sample data created.';
35

```

```
-- Stored procedures: Get paged customer orders and Add order
IF OBJECT_ID('dbo.usp_GetCustomerOrders', 'P') IS NOT NULL
    DROP PROCEDURE dbo.usp_GetCustomerOrders;
IF OBJECT_ID('dbo.usp_AddOrder', 'P') IS NOT NULL
    DROP PROCEDURE dbo.usp_AddOrder;

GO

CREATE PROCEDURE dbo.usp_GetCustomerOrders
    @CustomerId INT,
    @StartDate DATETIME = NULL,
    @EndDate DATETIME = NULL,
    @Page INT = 1,
    @PageSize INT = 10
AS
BEGIN
    SET NOCOUNT ON;
    IF @Page < 1 SET @Page = 1;
    DECLARE @Offset INT = (@Page - 1) * @PageSize;

    SELECT o.OrderId, o.CustomerId, o.OrderDate, o.Amount, o.Status,
    COUNT(*) OVER() AS TotalCount
    FROM dbo.Orders o
    WHERE o.CustomerId = @CustomerId
        AND (@StartDate IS NULL OR o.OrderDate >= @StartDate)
        AND (@EndDate IS NULL OR o.OrderDate <= @EndDate)
    ORDER BY o.OrderDate DESC
    OFFSET @Offset ROWS FETCH NEXT @PageSize ROWS ONLY;
END

GO

CREATE PROCEDURE dbo.usp_AddOrder
    @CustomerId INT,
    @OrderDate DATETIME,
    @Amount DECIMAL(10,2),
    @Status NVARCHAR(20),
```

```

    |   OFFSET @Offset ROWS FETCH NEXT @PageSize ROWS ONLY;
END

GO

CREATE PROCEDURE dbo.usp_AddOrder
    @CustomerId INT,
    @OrderDate DATETIME,
    @Amount DECIMAL(10,2),
    @Status NVARCHAR(20),
    @NewOrderId INT OUTPUT
AS
BEGIN
    SET NOCOUNT ON;
    BEGIN TRY
        BEGIN TRANSACTION;
        INSERT INTO dbo.Orders (CustomerId, OrderDate, Amount, status)
        VALUES (@CustomerId, @OrderDate, @Amount, @Status);
        SET @NewOrderId = CAST(SCOPE_IDENTITY() AS INT);
        COMMIT TRANSACTION;
    END TRY
    BEGIN CATCH
        IF XACT_STATE() <> 0 ROLLBACK TRANSACTION;
        THROW;
    END CATCH
END

GO

PRINT 'Stored procedures created.';
```

Output:

```
Falling back to mock demo output so you can see expected behavior.
Running mock demo (no DB connection)
Adding a new order for customer 1
New OrderId: 101

Fetching page 1 of orders for customer 1
Total orders for customer: 3
[{'OrderId': 101, 'CustomerId': 1, 'OrderDate': '2025-10-20', 'Amount': 199.99, 'status': 'Processing'}, {'OrderId': 3, 'CustomerId': 1, 'OrderDate': '2025-10-01', 'Amount': 250.0, 'status': 'Processing'}, {'OrderId': 2, 'CustomerId': 1, 'OrderDate': '2025-09-25', 'Amount': 45.0, 'status': 'Shipped'}]
(.venv) PS C:\Users\LENOVO\OneDrive\Desktop\lab exam 04 2.ox
```